

# THE NEURAL PERSPECTIVE

DEEP LEARNING SIMPLIFIED

HOME • TUTORIALS • READINGS • CONTACT • ABOUT

## LOGISTIC REGRESSION

*Edit*

### I. OBJECTIVE:

We use different equations depending on the number of output classes. With 2 classes, we will use binomial cross entropy for loss and more than 2 classes involves using the cross-entropy with softmax.

Regression	Binomial	Multinomial (n>2)
$\hat{y}$	$\hat{y} = \frac{1}{1 + e^{-\theta X}}$	$\hat{y} = \frac{e^{\theta_y X}}{\sum e^{\theta_i X}}$
cross entropy	$J(\theta) = -\frac{1}{n} \sum_i [y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)]$	$J(\theta) = -\sum_i y_i \ln(\hat{y}_i)$
forward pass	$[NX1] = [NXD][DX1]$	$[NXC] = [NXD][DXC]$
response	$y = 0 \text{ or } 1$	$y = \text{one-hot-encoded}$

### II. MULTINOMIAL CROSS ENTROPY:

Example of how to calculate the cross-entropy loss for a 3 class problem.

$$J(\theta) = -\sum_i y_i \ln(\hat{y}_i)$$

Ex:	Computed ( $\hat{y}$ )	Targets (y)
	[0.3, 0.3, 0.4]	[0, 0, 1]

$$J(\theta) = -\sum_i [0 * \ln(0.3) + 0 * \ln(0.3) + 1 * \ln(0.4)] = -\ln(0.4)$$

### III. BACKPROPAGATION:

With the multinomial cross entropy, you can see that we only keep the loss contribution from the correct class. Usually, with neural nets, this will be case if our outputs are sparse (just 1 true class). Therefore, we can rewrite our loss into just a sum(-log( $\hat{y}_i$ )) where  $\hat{y}_i$  will just be the probability of the correct class. We just replace  $y_{-i}$  (true y) with 1 and for the probabilities for the other classes, doesn't matter because their  $y_{-i}$  is 0. This is referred to as negative log likelihood.

**Note:** The  $y^i$  below is just the probability for the correct class, so the loss is only affected by the probability of the correct class. However, the gradients DO take into account the probabilities for the other classes. This allows us to change all the weights towards minimizing the loss function. Also all the ys below are  $\hat{y}$ s.

$$J(\theta) = -\sum_i \ln(\hat{y}_i)$$

Negative Log Likelihood

$$\frac{\partial J}{\partial W_j} = \frac{\partial J}{\partial y} \frac{\partial y}{\partial W_j} = -\frac{1}{y} \frac{\partial y}{\partial W_j} = -\frac{1}{\sum e^{W_j X}} \frac{\sum e^{XW} e^{W_j X} - e^{W_j X} e^{W_j X} X}{(\sum e^{XW})^2} = \frac{X e^{W_j X}}{\sum e^{XW}} = X P$$

$$\frac{\partial J}{\partial W_y} = \frac{\partial J}{\partial y} \frac{\partial y}{\partial W_y} = -\frac{1}{y} \frac{\partial y}{\partial W_y} = -\frac{1}{\sum e^{W_y X}} \frac{\sum e^{XW} e^{W_y X} - e^{W_y X} e^{W_y X} X}{(\sum e^{XW})^2} = \frac{1}{P} (X P - X P^2) = X (P - 1)$$

$$W_j = W_j - \alpha \frac{\partial J}{\partial W_j}$$

- SEARCH -

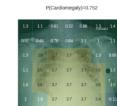
Search ...

- FOLLOW ME ON TWITTER -

Tweets by @GokuMohandas

 @GokuMohandas

Your model may be performing really well by incorrectly focusing on confounding features (extraneous influencers in the data that aren't accounted for). Check out [@johnzrech's post](#) where x-ray stickers unintentionally influenced the classifications: [medium.com/@jrzrech/what-a...](#)



Embed

View on Twitter

- RECENT POSTS -

Update on Embeddings (Spring 2017)

Exploring Sparsity in Recurrent Neural Networks

Question Answering from Unstructured Text by Retrieval and Comprehension

Overcoming Catastrophic Forgetting in Neural Networks

Opening the Black Box of Deep Neural Networks via Information

$$\partial W_i$$

## IV. SOFTMAX CLASSIFIER IMPLEMENTATION:

Naive Implementation:

```

1 | loss = 0.0
2 | dW = np.zeros_like(W)
3 | num_train = X.shape[0]
4 | num_classes = W.shape[1]
5 |
6 | for i in xrange(num_train):
7 |
8 |     scores = X[i, :].dot(W)
9 |     scores -= np.max(scores)
10 |    normalized_scores = np.exp(scores) / np.sum(np.exp(sco
11 |
12 |    for j in xrange(num_classes):
13 |
14 |        if j == y[i]:
15 |            loss += -np.log(normalized_scores[j])
16 |            dW[:, j] += (normalized_scores[j] - 1.0) * X[i]
17 |        else:
18 |            dW[:, j] += (normalized_scores[j] - 0.0) * X[i]
19 |
20 |    loss /= num_train
21 |    dW /= num_train
22 |    loss += 0.5 * reg * np.sum(W * W)
23 |    dW += reg * W

```

Vectorized Implementation:

```

1 | loss = 0.0
2 | dW = np.zeros_like(W)
3 | num_train = X.shape[0]
4 | num_classes = W.shape[1]
5 |
6 | scores = np.exp(X.dot(W))
7 | normalized_scores = (scores.T / np.sum(scores, axis=1)).T
8 |
9 | ground_truth = np.zeros_like(normalized_scores)
10 | ground_truth[range(num_train), y] = 1.0 # correct class
11 |
12 | loss = np.sum(np.sum(-np.log(normalized_scores[range(num_t
13 | dW = X.T.dot(normalized_scores - ground_truth)
14 |
15 | loss /= num_train
16 | dW /= num_train
17 | loss += 0.5 * reg * np.sum(W * W)
18 | dW += reg * W

```

## V. CODE BREAKDOWN:

We will be using `load_data()` to load the MNIST data if needed and separate into train and test sets.

```

1 | import tensorflow as tf
2 | import numpy as np
3 | import input_data
4 |
5 | class parameters():
6 |
7 |     def __init__(self):
8 |         self.LEARNING_RATE = 0.05
9 |         self.NUM_EPOCHS = 500
10 |        self.DISPLAY_STEP = 1 # epoch
11 |
12 |    def load_data():
13 |        mnist = input_data.read_data_sets("MNIST_data/", one_h
14 |        trainX, trainY, testX, testY = mnist.train.images, mni
15 |        return trainX, trainY, testX, testY

```

We will create a model and use softmax cross-entropy with logits as our loss function. And `step()` will do the training/validation steps for us. Note the `forward_only` argument. When we are validating on the test set, we do not want to train on this set, so we will not do any operations involving the optimizer.

```

1 | def create_model(sess, learning_rate):
2 |     tf_model = model(learning_rate)
3 |     sess.run(tf.initialize_all_variables())
4 |     return tf_model
5 |
6 | class model(object):
7 |
8 |     def __init__(self, learning_rate):
9 |
10 |         # Placeholders
11 |         self.X = tf.placeholder("float", [None, 784])
12 |         self.y = tf.placeholder("float", [None, 10])
13 |
14 |         # Weights
15 |         with tf.variable_scope('weights'):
16 |             W = tf.Variable(tf.random_normal([784, 10]), st
17 |
18 |             self.logits = tf.matmul(self.X, W)
19 |             self.cost = tf.reduce_mean(tf.nn.softmax_cross_ent
20 |             self.optimizer = tf.train.GradientDescentOptimizer
21 |
22 |             # Prediction
23 |             self.prediction = tf.argmax(self.logits, 1)
24 |

```

```

25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40

```

We will train the model using the entire train/test batches at each step.

```

1 def train(FLAGS):
2
3     with tf.Session() as sess:
4
5         model = create_model(sess, FLAGS.LEARNING_RATE)
6         trainX, trainY, testX, testY = load_data()
7
8         for epoch_num in range(FLAGS.NUM_EPOCHS):
9             prediction, training_loss, _ = model.step(sess
10
11             # Display
12             if epoch_num*FLAGS.DISPLAY_STEP == 0:
13                 print "EPOCH %i: \n Training loss: %.3f, T
14                 epoch_num, training_loss, model.step(s
15
16         if __name__ == '__main__':
17             FLAGS = parameters()
18             train(FLAGS)

```

## VI. RAW CODE:

[GitHub Repo](#) (Updating all repos, will be back up soon!)



Posted in: Regression

Tagged: Tutorials

← Linear Regression

Vanilla Neural Network →

## 3 THOUGHTS ON “LOGISTIC REGRESSION”

FIDELA January 30, 2018 at 5:41 pm

[EDIT](#) [REPLY→](#)



\_\_\_123\_\_\_Logistic Regression – The Neural Perspectives\_\_\_123\_\_\_

★ Like

FRIGG February 4, 2018 at 3:45 pm

[EDIT](#) [REPLY→](#)



\_\_\_123\_\_\_Logistic Regression – The Neural Perspectives\_\_\_123\_\_\_

★ Like

KATIA BUSER February 5, 2018 at 7:52 pm

[EDIT](#) [REPLY→](#)



They would not assume it's joyous if it concerned a planet  
burning and all of us getting stumped flat by a planet  
now, would they? We consider the number '5' as making use of to acceptable groups  
of any entities in anyway-to five fishes, 5 kids, 5  
apples, five days <https://math-problem-solver.com/>.  
60 minute response time or its free.

★ Like

## LEAVE A REPLY

Enter your comment here...

